

# Neural Nets for text classification

Alexandre Allauzen

Winter 2024



# Roadmap

Introduction

Word embeddings

Convolution for text

References

# Outline

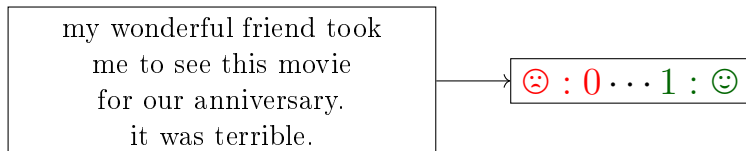
Introduction

Word embeddings

Convolution for text

References

# Text classification/rating



How to represent the input text ?

- Bag of features (words, ... )
- Really represent the sentence as a whole

# Bag of words (BOW)

*this movie is just great , with a great music , while a bit long*

vocabulary	binary bag	count bag	tf.idf bag	...
awesome	0	0	0	...
great	1	2	1.9	...
long	1	1	2.5	...
the	0	0	0	...
this	1	1	0.1	...

A basic vectorial representation of text

$$\mathbf{x} = \begin{pmatrix} 0 \\ 2 \\ 1 \\ 0 \\ 1 \end{pmatrix} \in \mathbb{R}^D$$

$$\left. \begin{array}{l} \textit{awesome} \\ \textit{great} \\ \textit{long} \\ \textit{the} \\ \textit{this} \end{array} \right\} D$$

# A simple problem

## Assumptions

- Let define a finite set of known words: the vocabulary  $\mathcal{V}$
- A text is a vector  $\mathbf{x}$  of dimension  $D = |\mathcal{V}|$
- Each component encodes the presence of a word

## Then machine learning

- Naive Bayes
- SVM, Random Forrest, ...
- Logistic Regression



# Logistic regression

The class  $c$  is the outcome of the binary random variable  $C$

The sigmoid/logistic function

$$a = w_0 + \mathbf{w}^t \mathbf{x} \in \mathbb{R}$$

$$\sigma(a) = \frac{e^a}{1 + e^a} = \frac{1}{1 + e^{-a}} \text{ and } y = P(C = 1|\mathbf{x}) = \sigma(w_0 + \mathbf{w}^t \mathbf{x})$$



# Training a Logistic regression model

- The parameters are  $\boldsymbol{\theta} = (w_0, \mathbf{w})$ ,
- The i.i.d dataset:  $\mathcal{D} = (\mathbf{x}_{(i)}, c_{(i)})_{i=1}^n$

## Loss function minimization

$$\begin{aligned}\mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) &= - \sum_{i=1}^n \log(P(C = c_{(i)} | \mathbf{x}; \boldsymbol{\theta})) \\ &= - \sum_{i=1}^n (c_{(i)} \log(y_{(i)}) + (1 - c_{(i)}) \log(1 - y_{(i)})) \\ y_{(i)} &= \sigma(w_0 + \mathbf{w}^t \mathbf{x}_{(i)})\end{aligned}$$

## Optimization method

Stochastic Gradient Descent, or improved version (ADAM, L-BFGS, ...)



# Outline

Introduction

Word embeddings

Convolution for text

References

## Back to logistic regression

$$\mathbf{x} = \begin{pmatrix} 0 \\ 2 \\ 1 \\ 0 \\ 1 \end{pmatrix} \in \mathbb{R}^D \quad \left. \begin{array}{l} \textit{awesome} \\ \textit{great} \\ \textit{long} \\ \textit{the} \\ \textit{this} \end{array} \right\} D$$

For one input text:


$$w_0 + \mathbf{w}^t \mathbf{x} = w_0 + \textcolor{red}{2} \times w_2 + w_3 + w_5$$

The class is positive ( $y = 1$ ) if

$$\begin{aligned} w_0 + 2 \times w_2 + w_3 + w_5 &> 0 \\ 2 \times w_{\textit{great}} + w_{\textit{long}} + w_{\textit{this}} + &> -w_0 \end{aligned}$$

# A limited representation of words

With the logistic regression model on a bag of words:



Consider the two following examples:

the end is <b>really bad</b>		☹️ $\Rightarrow w_{\text{bad}}$	$\searrow$
the <b>bad</b> guy is <i>awesome</i>		😊 $\Rightarrow w_{\text{bad}}$	$\searrow, w_{\text{awesome}} \nearrow$

Multiple dimensions could help to:

- represent different usage
- consider the context
- leverage more from sparse, sometime ambiguous observations.

# A simple model for document classification - part 1

## Idea

- The word representation could be shared among classes
- While their interpretation depends on the class

## Input representation and composition

$$\mathbf{R} \times \mathbf{x} = \begin{pmatrix} \vdots & \vdots & \vdots & \vdots & \vdots \\ \mathbf{v}_1 & \mathbf{v}_2 & \mathbf{v}_3 & \mathbf{v}_4 & \mathbf{v}_5 \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{pmatrix} \times \begin{pmatrix} 0 \\ \mathbf{2} \\ \mathbf{1} \\ 0 \\ \mathbf{1} \end{pmatrix} = 2 \times \mathbf{v}_2 + \mathbf{v}_3 + \mathbf{v}_5 = \mathbf{d}$$

# A simple model for document classification - part 2

## Classification

$$\begin{aligned}P(y|\mathbf{x}) &= \text{softmax}(\mathbf{W}^o \mathbf{d}) = \text{softmax}(\mathbf{W}^o \times \mathbf{R}\mathbf{x}), \text{ or} \\&= \text{softmax}(\mathbf{W}^o \times f(\mathbf{R}\mathbf{x})),\end{aligned}$$

with  $f$  a non-linear activation function.

## Parameters


$$\theta = (\mathbf{R}, \mathbf{W}^o) \rightarrow \text{to learn !!}$$

## Reminder

If  $\mathbf{y} = \text{softmax}(\mathbf{a})$ ,  $\mathbf{y}$  is a vector and  $\mathbf{a}$  is called the logit vector

$$y_i = \frac{e^{a_i}}{\sum_j e^{a_j}}$$

# A first neural network



- $\mathbf{x} : (|\mathcal{V}|, 1)$
- $\mathbf{R} : (K, |\mathcal{V}|)$
- $\mathbf{d} : (K, 1)$
- $\mathbf{W}^o : (1, K)$
- $y : (1, 1)$

$$\mathbf{d} = \mathbf{R} \times \mathbf{x}$$

$$y = \sigma(\mathbf{W}^o \times \mathbf{d})$$

# Word embeddings

## Definitions:

- To each word, a continuous vector is associated: its **embedding**.
- The matrix  $\mathbf{R}$  is called the **look-up table** and store the word embeddings.

## Note:

- The term *look-up* comes from the real operation  $\mathbf{R} \times \mathbf{x}$  is only theoretical !
- No computational cost, only storage and trainability challenge (enough observations for each word, Zipf, ...)
- **Pre-training** and **fine-tuning**

# Unsupervised Pre-training of Word Embeddings

## The question

- How to efficiently learn word representation
- based on the observation of raw texts ?

## Distributional representations

*You shall know a word by the company it keeps (Firth, J. R., 1957)*

and

*Words are similar if they appear in similar contexts (Harris 1954).*

## In practice

Word2Vec [6]



# Context Bag of Words (CBOW)

The game

southern trees [??] strange fruits


Guess the word in the middle !

# Context Bag of Words (CBOW)

The game

southern trees [??] strange fruits

Guess the word in the middle !



# CBOW: details

## Fast pre-training of word embeddings

- Introduced in [6] as a simplification of [1] (neural language model)
- Trained with negative sampling (Closed to Noise Contrastive Estimation [2])
- An efficient and tractable approximation of the count based method [5]

## Other flavor

- Skip-gram [6]
- Glove [7]
- Fasttext [3]

# CBOW: Maximum Likelihood Estimate

In  $P(w|\mathbf{x}; \boldsymbol{\theta})$ :

- predict the word  $w$  in the middle,
- given  $\mathbf{x}$  the context.

MLE

$$\mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) = - \sum_{i=1}^n \log(P(C = w|\mathbf{x}; \boldsymbol{\theta})),$$

- The probability distribution over  $\mathcal{V}$  is given by a softmax
- The set of possible outcomes is  $\mathcal{V}$ .

## Cost of the softmax

$$\mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) = - \sum_{(\mathbf{x}, \hat{w}) \in \mathcal{D}} \log P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})$$

$$P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x}) = \frac{e^{s_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})}}{\sum_{w' \in \mathcal{V}} e^{s_{\boldsymbol{\theta}}(w'|\mathbf{x})}}$$

$$\log P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x}) = s_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x}) - \log \left( \sum_{w' \in \mathcal{V}} e^{s_{\boldsymbol{\theta}}(w'|\mathbf{x})} \right)$$

$$\frac{\partial \log P_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})}{\partial \boldsymbol{\theta}} = \frac{\partial s_{\boldsymbol{\theta}}(\hat{w}|\mathbf{x})}{\partial \boldsymbol{\theta}} - \underbrace{\sum_{w' \in \mathcal{V}} P_{\boldsymbol{\theta}}(w'|\mathbf{x}) \frac{\partial s_{\boldsymbol{\theta}}(w', \mathbf{x})}{\partial \boldsymbol{\theta}}}_{\text{costly!}}$$

# Negative sampling

Recast the problem as a binary classification task:

- Positive examples:  $(\mathbf{x}, w) \in \mathcal{D}$
- Negative examples:  $(\mathbf{x}, \tilde{w})$ , with  $\tilde{w} \sim \mathcal{V}$

Use a binary classifier !

In practice:

- for one positive example  $\sim \mathcal{D}$
- sample  $K$  negative and random samples from  $\mathcal{V}$
- $K$  is small (compared to the size of  $\mathcal{V}$ )
- the noise distribution does matter

# Outline

Introduction

Word embeddings

Convolution for text

References

# Local contexts

the	end	is	very	bad	but	what	a	great	music
-----	-----	----	------	-----	-----	------	---	-------	-------



# Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ $very \rightarrow bad ++$						

## Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ $very \rightarrow bad++$						
			$\underbrace{\hspace{2.5cm}}$ $but \text{ will change } bad$						

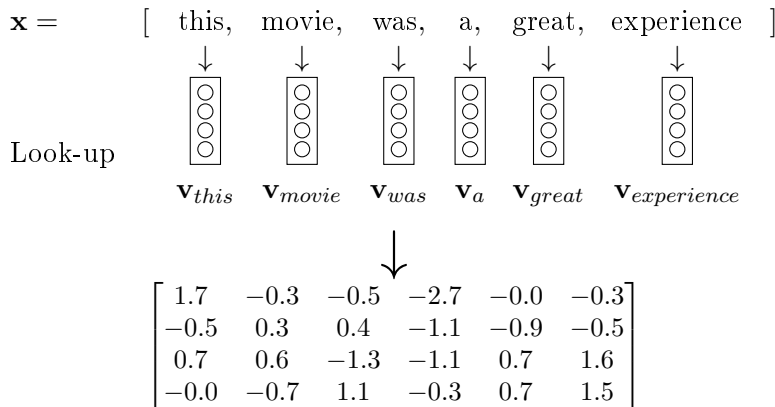
# Local contexts

the	end	is	very	bad	but	what	a	great	music
			$\underbrace{\hspace{1.5cm}}$ <i>very</i> $\rightarrow$ <i>bad</i> ++						
			$\underbrace{\hspace{2.5cm}}$ <i>but</i> will change <i>bad</i>						
		$\underbrace{\hspace{3.5cm}}$ <i>bad</i> is for <i>end</i> not <i>music</i>						$\underbrace{\hspace{1.5cm}}$ <i>great</i> is for <i>music</i> not fo <i>end</i>	

## Motivations

- Local contextualisation
- Global view of the sentence

## Another view of a sentence



## Convolutional Neural Networks for Sentence Classification

- A short paper of 2014 [4]
- A simple and SOTA paper on text classification

# Convolution in 1D

Extract a frame, or a window, and apply a “filter”

The filter

Kernel size of  $ks = 2$

$w_{1,1}$	$w_{1,2}$
$w_{2,1}$	$w_{2,2}$
$w_{3,1}$	$w_{3,2}$
$w_{4,1}$	$w_{4,2}$



The input sequence

$L = 6$  vectors in  $\mathbb{R}^D$ ,  $D = 4$

$x_{1,1}$	$x_{1,2}$	$x_{1,3}$	$x_{1,4}$	$x_{1,5}$	$x_{1,6}$
$x_{2,1}$	$x_{2,2}$	$x_{2,3}$	$x_{2,4}$	$x_{2,5}$	$x_{2,6}$
$x_{3,1}$	$x_{3,2}$	$x_{3,3}$	$x_{3,4}$	$x_{3,5}$	$x_{3,6}$
$x_{4,1}$	$x_{4,2}$	$x_{4,3}$	$x_{4,4}$	$x_{4,5}$	$x_{4,6}$

The output value (output channel)

$$\text{At time } t = 1, \text{ } h_1 = \sum_{i,j} w_{i,j} \times x_{i,j}$$

# Convolution in 1D : time $t = 2$

Stride = 1

The filter

Kernel size of  $ks = 2$

$w_{1,1}$	$w_{1,2}$
$w_{2,1}$	$w_{2,2}$
$w_{3,1}$	$w_{3,2}$
$w_{4,1}$	$w_{4,2}$



The input sequence

$L = 6$  vectors in  $\mathbb{R}^D$ ,  $D = 4$

$x_{1,1}$	$x_{1,2}$	$x_{1,3}$	$x_{1,4}$	$x_{1,5}$	$x_{1,6}$
$x_{2,1}$	$x_{2,2}$	$x_{2,3}$	$x_{2,4}$	$x_{2,5}$	$x_{2,6}$
$x_{3,1}$	$x_{3,2}$	$x_{3,3}$	$x_{3,4}$	$x_{3,5}$	$x_{3,6}$
$x_{4,1}$	$x_{4,2}$	$x_{4,3}$	$x_{4,4}$	$x_{4,5}$	$x_{4,6}$

The output value (output channel)

$$\text{At time } t = 2, \text{ } h_2 = \sum_{i,j} w_{i,j} \times x_{i+1,j}$$

# Convolution in 1D : time $t = 5$

Stride = 1

The filter

Kernel size of  $ks = 2$

The input sequence

$L = 6$  vectors in  $\mathbb{R}^D$ ,  $D = 4$

$w_{1,1}$	$w_{1,2}$
$w_{2,1}$	$w_{2,2}$
$w_{3,1}$	$w_{3,2}$
$w_{4,1}$	$w_{4,2}$




$x_{1,1}$	$x_{1,2}$	$x_{1,3}$	$x_{1,4}$	$x_{1,5}$	$x_{1,6}$
$x_{2,1}$	$x_{2,2}$	$x_{2,3}$	$x_{2,4}$	$x_{2,5}$	$x_{2,6}$
$x_{3,1}$	$x_{3,2}$	$x_{3,3}$	$x_{3,4}$	$x_{3,5}$	$x_{3,6}$
$x_{4,1}$	$x_{4,2}$	$x_{4,3}$	$x_{4,4}$	$x_{4,5}$	$x_{4,6}$

The output value (output channel)

$$\text{At time } t = 5, \text{ } h_5 = \sum_{i,j} w_{i,j} \times x_{i+5,j}$$

# Feature extraction for subsequences

## Embeddings



## After the convolution

$$\mathbf{h} = ( \underbrace{h_1}_{(\text{this,movie})}, \underbrace{h_2}_{(\text{movie,was})}, \underbrace{h_3}_{(\text{was,a})}, \underbrace{h_4}_{(\text{a,great})}, \underbrace{h_5}_{(\text{great,experience})} )$$



# Convolution 1D on a embedding sequence

## Local features extraction

- Each Kernel / Filter application gives **one local feature**.
- The feature relates to a  $n$ -gram.
- The **kernel size** defines the scope of each feature.
- The **stride** controls the amount of slides.

A sequence of vector  $\rightarrow$  a sequence of features


## More features ?

- More Filters !
- More "output channels"

# Convolution with two output channels

Output channels = 2

Input channels = 1




The output value (output channel)

$$h_{1,1} = \sum_{i,j} w_{i,j}^{(1)} \times x_{i,j}$$

$$h_{2,1} = \sum_{i,j} w_{i,j}^{(2)} \times x_{i,j}$$


## Another view for two output channels



- Two filters applied to the same frame (or window)
- Two projections
- $\mathbf{W}$ : the parameters of the filters
- $\mathbf{W}$  is learnt

# Pool !

Compress the "local" information along one dimension (e.g time)




- Mean pooling
- Max pooling, and  $k$ -max pooling

For example with:

$$\begin{bmatrix} 1.7 & -0.3 & -0.5 & -2.7 & -0.0 & -0.3 \\ -0.5 & 0.3 & 0.4 & -1.1 & -0.9 & -0.5 \end{bmatrix} \rightarrow \begin{bmatrix} -0.3 \\ -0.4 \end{bmatrix} \text{ or } \begin{bmatrix} 1.7 \\ 0.4 \end{bmatrix} \text{ or } \dots$$

Pooling can also apply to sliding windows

# More Convolutions and pooling



And they can be combined (concatenation).

# Convolutional Neural Networks for Sentence Classification

## A summary of [4]

- Window (kernel) sizes : 3, 4, 5 with 100 feature maps for each
- Static/non-static/random/multi-channel word embeddings
- Auxiliary data for word embeddings:  $\tilde{w}2v$  trained on 100 billion words from Google News ( $dim = 300$ )
- dropout on the penultimate layer (after the max-pooling)
- Relu and early stopping

# CNN applications for NLP

## Word level

- The unit = a word (the vocabulary ?)
- Compose word representation to derive a sentence representation
- Extract  $n$ -gram patterns (phrasal)

## Char level

- The unit = the char (closed vocab)
- Compose chars to infer a word representation
- Extract morphological features (mostly concatenative)  
unbelievable, untractable, believer, writter, forever, ...

# Outline

Introduction

Word embeddings

Convolution for text

References



- [1] Yoshua Bengio and Réjean Ducharme. “A Neural Probabilistic Language Model”. In: *Advances in Neural Information Processing Systems (NIPS)*. Vol. 13. Morgan Kaufmann, 2001.
- [2] M. Gutmann and A. Hyvärinen. “Noise-contrastive estimation: A new estimation principle for unnormalized statistical models”. In: *Proceedings of the International Conference on Artificial Intelligence and Statistics (AISTATS)*. Ed. by Y.W. Teh and M. Titterton. Vol. 9. 2010, pp. 297–304.
- [3] Armand Joulin et al. “Bag of Tricks for Efficient Text Classification”. In: *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 2, Short Papers*. Valencia, Spain: Association for Computational Linguistics, Apr. 2017, pp. 427–431. URL: <https://www.aclweb.org/anthology/E17-2068>.
- [4] Yoon Kim. “Convolutional Neural Networks for Sentence Classification”. In: *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Doha, Qatar: Association for Computational Linguistics, 2014, pp. 1746–1751. URL: <http://www.aclweb.org/anthology/D14-1181>.
- [5] Oren Melamud, Ido Dagan, and Jacob Goldberger. “PMI Matrix Approximations with Applications to Neural Language Modeling”. In: *CoRR* abs/1609.01235 (2016). arXiv: 1609.01235. URL: <http://arxiv.org/abs/1609.01235>.

- [6] Tomas Mikolov et al. “Distributed Representations of Words and Phrases and their Compositionality”. In: *Advances in Neural Information Processing Systems 26*. Ed. by C.J.C. Burges et al. Curran Associates, Inc., 2013, pp. 3111–3119. URL: <http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf>.
- [7] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. “GloVe: Global Vectors for Word Representation”. In: *Empirical Methods in Natural Language Processing (EMNLP)*. 2014, pp. 1532–1543. URL: <http://www.aclweb.org/anthology/D14-1162>.